SEVARI

Autonomous Business Operations Engine

Seven specialist AI agents. One unified brain. Your business runs itself — with you in control of every decision that matters. Built on Serenity's battle-tested architecture. Deployed globally from day one.

Design Brief v1.0 | April 2026
Classification: CONFIDENTIAL — INTERNAL USE
Architecture: Built on Operion Serenity DNA
Platform: IONOS VPS | 8 vCore | 16GB RAM | 480GB NVMe SSD
Infrastructure: n8n (self-hosted) | Ollama (CPU) | PostgreSQL
Author: Dan Hadfield | Sole Director, Operion Automation Systems Ltd

SECTION 01

What Sevari Does

Sevari does everything a team of employees would do to build and run a business — code it, market it, sell it, support it, invoice it, monitor it — autonomously, every day, using specialist AI agents with defined boundaries and human approval at every critical point.

> **IN PLAIN ENGLISH**

> You tell Sevari what your business is. Sevari writes your code and pushes it to GitHub. Builds and deploys your website. Writes your marketing content. Creates and sends your outreach emails. Handles your customer support. Generates your invoices. Tracks your revenue. Monitors your infrastructure. Analyses your market. Plans tomorrow's work. Sends you a morning brief and an evening summary. You approve what matters from your phone. Everything else runs without you.

Sevari vs Competitors

DIMENSION	COMPETITORS (POLSIA AI ETC.)	SEVARI
Agent model	One generalist AI loop does everything	Seven specialist agents, each with dedicated models, tools, and boundaries. HIVE Mind coordinates.
Confirmation gates	None or minimal. AI acts autonomously.	Guardian Layer. Three gate levels. Financial actions always human-approved.
Audit trail	Dashboard. No immutability. No replay.	Immutable append-only log. SHA-256 chain hashing. Evidence-linked morning briefs.
Failure handling	Retry. No failsafe chain.	Six-tier failsafe. Model cascade → templates (zero AI, cannot fail) → human escalation.
Data sovereignty	Their cloud. Your business dies if they die.	Your VPS. UK datacenter. Full export. Zero vendor lock-in.
Business memory	None. Each cycle starts fresh.	HIVE Mind learns from every operation and every approve/reject decision you make.
Revenue share	Up to 20% of your revenue.	Fixed monthly subscription. £79–£399. Zero revenue share.

Data protection	Varies. Often US-only compliance.	GDPR, UK GDPR, CCPA, PIPEDA, LGPD. AES-256-GCM PII encryption. Immutable audit.

SECTION 02

How Sevari Works

Three interaction modes. One architecture. Everything Guardian-gated.

Mode 1: Command

You send a task via dashboard, API, email, or webhook. Sevari routes it to the right specialist agent, executes it, Guardian-gates the output, and either auto-delivers or queues for your approval depending on the gate level.

Mode 2: Approve

Morning brief arrives at 07:00. Here's what I did overnight. Here's what needs your sign-off. Tap approve or reject from your phone. HIVE Mind learns from every decision.

Mode 3: Autonomous

Scheduled agents run daily. HERALD writes content. COMPASS scans the market. LEDGER reconciles. SENTINEL checks health. You only get pulled in when Guardian says so.

> **COMMUNICATION MODEL — INHERITED FROM SERENITY**
>
> Sevari broadcasts OUT only (one-way transmitter). HIVE Mind, AI Ecosystem, and Knowledge Platform communicate bidirectionally with each other. Nothing flows back to Sevari's core without Dan's explicit approval. HIVE Mind learns from Dan's approve/reject decisions. No system can change Sevari without Dan's sign-off.

Task Lifecycle

STEP	WHAT HAPPENS	N8N ZONE
1. Intake	Task arrives via API, form, webhook, email, or schedule trigger. Tagged with source, timestamp, IP.	INTAKE (shared with Serenity)
2. Guardian	8 threat categories scanned. Phi AI classification for edge cases. Nothing enters core without clearance.	GUARDIAN (shared)
3. Credentials	API key verified via PostgreSQL (SHA-256 hashed). Business status + tier limits enforced.	CREDENTIALS (shared)
4. Agent Router	Task classified to target agent. Confidence scored. Fallback to BASTION if ambiguous.	SEVARI ROUTER (new)
5. Agent Execution	Specialist agent runs with dedicated prompt, model, and tool access. Six-tier failsafe chain.	SEVARI AGENTS (new)
6. Output Guard	PII scrub. Identity leak prevention. Cross-tenant isolation check.	OUTPUT (shared)
7. Guardian Gate	Output classified by risk. Auto-approve, queue-and-notify, or require-approval.	GUARDIAN GATE (new)
8. HIVE Mind	Pattern recorded. Recommendations generated. Learning logged. Approval queue updated.	HIVE MIND (shared)
9. Audit	Immutable trail. SHA-256 chain hash. Usage metrics updated. PostgreSQL writes.	AUDIT (shared)

SECTION 03

The Seven Agents

Each agent has a defined role, dedicated prompt, specific tool access, and Guardian gate level. No agent can do another agent's job. No agent can bypass its own boundaries. HIVE Mind coordinates them.

FORGE

Code & Development

Generates application code, fixes bugs, writes tests, refactors, handles dependency updates. Uses the strongest available code model.

Tools: GitHub API, local IDE, test runners, linters | **Model:** qwen2.5-coder:7b | **Gate:** Staging auto-approve, production require-approval

HERALD

Marketing & Content

Writes blog posts, LinkedIn content, tweets, newsletters, case studies. Maintains brand voice. Runs content calendar.

Tools: CMS, email platforms, social APIs | **Model:** mistral:7b | **Gate:** All content require-approval before publication

VANGUARD

Sales & Outreach

Drafts outreach emails, manages lead pipeline, tracks prospects, qualifies leads, schedules follow-ups.

Tools: CRM, email, scheduling APIs | **Model:** mistral:7b | **Gate:** All external communications require-approval

BASTION

Customer Support

Handles inbound support queries, routes issues, drafts responses, tracks resolution. This is Serenity's enquiry engine running inside Sevari — identical pipeline, same failsafe chain.

Tools: Email, webhook, knowledge base | **Model:** mistral:7b (failsafe: llama3.2:3b → gemma2:2b → phi3:mini → templates) | **Gate:** Auto-response auto-approve, AI responses queue-and-notify

LEDGER

Finance & Invoicing

Generates invoices, tracks revenue, reconciles payments, produces financial summaries. Every financial action requires human approval — no exceptions.

Tools: Stripe, accounting APIs | **Model:** phi3:mini (precision over creativity) | **Gate:** ALWAYS REQUIRE-APPROVAL

COMPASS

Market Intelligence

Scans competitor activity, analyses market trends, identifies opportunities, produces weekly intelligence reports.

Tools: Web scraping, RSS, analytics APIs | **Model:** mistral:7b | **Gate:** Reports auto-approve to dashboard, strategic recommendations require-approval

SENTINEL

Infrastructure & Monitoring

Monitors system health, checks uptime, tracks resource usage, alerts on anomalies, produces operational reports.

Tools: System APIs, health checks, log analysis | **Model:** phi3:mini (fast, lightweight) | **Gate:** Health reports auto-approve, infrastructure changes require-approval

SECTION 04

User Interaction Model

Task Submission

Users submit tasks the same way businesses submit enquiries to Serenity — via API, web form, webhook, or email. The task format extends the enquiry format with a `target_agent` field and `task_type` classification.

```
{
  "business_id": "BIZ-USER-001",
  "api_key": "IK-...",
  "subject": "Write a blog post about our new product launch",
  "body": "We're launching ProductX next Tuesday. Key features: A, B, C. Ta
  "target_agent": "HERALD",
  "task_type": "content_creation",
  "priority": "standard"
}
```

If `target_agent` is omitted, Sevari's router classifies the task automatically using keyword analysis and AI classification — the same way Serenity classifies enquiry categories.

Morning Brief (07:00 daily)

Email digest sent to the business owner containing: overnight agent activity summary, items requiring approval (with approve/reject links), HIVE Mind recommendations, system health status, and key metrics.

Evening Summary (19:00 daily)

What was accomplished today. What's queued for tomorrow. Approval outcomes and HIVE learning. Revenue and pipeline updates from LEDGER and VANGUARD.

Approval Flow

GATE LEVEL	BEHAVIOUR	USED FOR
AUTO-APPROVE	Execute immediately. Log to audit.	Health checks, internal reports, template responses, data collection
QUEUE-AND-NOTIFY	Execute but hold delivery. Notify owner. Auto-approve if no response within timeout (configurable, default 2h).	AI-generated customer responses, content drafts, non-financial outreach
REQUIRE-APPROVAL	Hold. Notify owner. Do not execute until explicit approval. No timeout.	Financial transactions, production deployments,

		external communications, config changes

SECTION 05

Serenity DNA — What's Inherited

Sevari is not a new product. She is Serenity's architecture deployed first, with seven specialist agents added on top. Every component below is live, tested, and carried forward unchanged.

COMPONENT	SERENITY V6.0	SEVARI EXTENSION
Intake	5-channel webhook merge + normalise + priority queue	+ task format with target_agent field
Guardian	8 threat types, Phi AI classification, regex + AI dual scan	Unchanged — tasks scanned identically to enquiries
PII Protection	AES-256-GCM field encryption, PII scrubbing before AI, response guard	Unchanged
Credentials	SHA-256 hashed API keys, PostgreSQL verification, tier enforcement	Unchanged
Classification	10-category DB-first with fallback	+ 7-agent routing with confidence scoring
RAG	nomic-embed-text via Ollama, PostgreSQL knowledge base	Unchanged — all agents share the knowledge base
AI Failsafe	Mistral → Llama2 → Gemma → Phi →	Updated models for VPS: Mistral:7b → Llama3.2:3b → Gemma2:2b → Phi3:mini

	Templates → Human	→ Templates → Human
HIVE Mind	Dual-purpose: Serenity improvement + business intelligence. Dan-only approval.	Unchanged — starts learning from Sevari operations. When Serenity launches customer-facing, HIVE Mind is already trained.
Audit	Immutable, SHA-256 chain-hashed, GDPR-compliant	+ source_product field: 'sevari' or 'serenity'
Marketing	Daily self-marketing pipeline (LinkedIn + tweet), Guardian-gated	HERALD agent absorbs and extends this pipeline
Scheduling	Health (5min), HIVE (10min), Ecosystem (6hr), Maintenance (midnight), Email (2min)	+ Agent schedules: HERALD (daily 06:00), COMPASS (daily 05:00), LEDGER (weekly Mon 03:00), SENTINEL (hourly), Brief (07:00 + 19:00)
Operations	Init, Demo, Health, Selftest, Stripe, Onboard, Dashboard, GDPR, Guardian approve/reject	+ Task submission, Agent status, Brief history, Approval queue API
PostgreSQL	33 tables, 6 immutable audit tables	+ 5 tables: sevari_tasks, sevari_agent_runs, sevari_briefs, sevari_approvals, sevari_agent_config

> **CRITICAL ARCHITECTURE PRINCIPLE**
>
> HIVE Mind is Serenity's brain. Sevari launches first and trains HIVE Mind with real-world operations data. When Serenity launches customer-facing, she inherits months of learned patterns, prompt refinements, knowledge base entries, and approval history. Serenity doesn't launch cold. She launches with a fully trained intelligence layer.

SECTION 06

AI Models — Free, Local, VPS-Optimised

All models run via Ollama on the VPS. CPU inference only (no GPU). All models are free, open-weight, and commercially licensed.

MODEL	SIZE	RAM	ROLE IN SEVARI	LICENCE
`mistral:7b-q4`	4.1GB	~5GB	Primary agent model (HERALD, VANGUARD, BASTION, COMPASS). Tier 1 failsafe.	Apache 2.0
`llama3.2:3b`	2.0GB	~3GB	Tier 2 failsafe. Fast fallback.	Llama 3.2 Community
`gemma2:2b`	1.6GB	~2.5GB	Tier 3 failsafe. Lightweight tasks.	Gemma Terms of Use
`phi3:mini`	2.3GB	~3GB	Guardian scanning, LEDGER (precision), SENTINEL (fast). Tier 4 failsafe.	MIT
`qwen2.5-coder:7b-q4`	4.4GB	~5GB	FORGE agent. Code specialist.	Apache 2.0
`nomic-embed-text`	274MB	~500MB	RAG embeddings. Knowledge base search.	Apache 2.0

> **VPS MEMORY MANAGEMENT**
>
> 16GB RAM. Ollama loads one model at a time by default (auto-unloads after 5 minutes idle). PostgreSQL uses ~1-2GB. n8n uses ~500MB. OS overhead ~1GB. Effective inference budget: ~12GB. All models fit individually. Two concurrent models possible with the smaller ones. The failsafe chain ensures that if a model fails to load, the next one tries — all the way down to templates (zero AI, zero RAM, cannot fail).

Six-Tier Failsafe Chain

TIER	MODEL	TIMEOUT	BEHAVIOUR
1	mistral:7b-q4	90s	Primary. Full capability.
2	llama3.2:3b	60s	Reduced context but functional.
3	gemma2:2b	45s	Lightweight. Basic tasks.
4	phi3:mini	30s	Minimal. Classification and short outputs.
5	Templates	Instant	Zero AI. Deterministic. Cannot fail. Structured templates per agent and task type.
6	Human Escalation	N/A	Queued to owner with full context. HIVE Mind records the failure pattern.

SECTION 07

Global Data Protection — Legally Defensible

Sevari is designed for global compliance from day one. Data residency is UK (IONOS UK datacenter). The architecture provides the technical controls required by all major data protection frameworks.

Framework Coverage

FRAMEWORK	JURISDICTION	SEVARI COMPLIANCE
GDPR	EU/EEA	Art 5 (minimisation), Art 6 (lawful basis), Art 13/14 (transparency), Art 17 (erasure), Art 20 (portability), Art 25 (by design), Art 28 (processor agreements), Art 30 (records of processing), Art 32 (security), Art 33/34 (breach notification)
UK GDPR	United Kingdom	Identical to GDPR with UK ICO as supervisory authority. UK adequacy decision provides data flow to EU.

CCPA/CPRA	California, USA	Right to know, right to delete, right to opt-out of sale. Sevari never sells data. Erasure endpoint covers deletion. Privacy notice required.
PIPEDA	Canada	10 fair information principles. Consent, purpose limitation, retention, accuracy, openness, individual access. All covered by existing architecture.
LGPD	Brazil	Modelled on GDPR. Lawful basis, purpose limitation, data minimisation, transparency, security, accountability. Architecture covers all requirements.

Technical Controls

CONTROL	IMPLEMENTATION
PII Encryption	AES-256-GCM field-level encryption on all personal data fields. Key derived via SHA-256 from environment variable.
PII Scrubbing	NI numbers, card numbers, bank details (sort code + account) masked before any AI model sees the data.
API Key Security	SHA-256 hashed. Never stored raw. Verified against PostgreSQL hashes.
Response Guard	Output scanned for PII leaks, identity leaks (Sevari/Operion/AI mentions), and cross-tenant data leaks before delivery.
Immutable Audit	audit_trail, ai_decision_log, system_events tables: NO UPDATE, NO DELETE rules. SHA-256 chain hashing.
Erasure (Art 17)	POST endpoint. Hashes email for lookup. Cascading deletion across all non-audit tables. 30-day SLA.
Export (Art 20)	POST endpoint. JSON format. All data associated with the requesting individual. 30-day SLA.
Data Residency	IONOS UK datacenter. Data never leaves UK infrastructure. All AI inference runs locally on the VPS.
Retention	Configurable per tier: Starter 90 days, Growth 1 year, Scale 2 years,

	Enterprise custom.
Multi-Tenant Isolation	All queries scoped by business_id. API key verification required. No cross-tenant data access possible.

> **LEGAL DOCUMENTATION REQUIRED**
>
> Terms of Service, Privacy Policy, AI Disclaimer, GDPR Policy (Sevari as data processor, client as data controller), Cookie Policy, Acceptable Use Policy, Data Processing Agreement (DPA) template for Enterprise clients. All existing from the Operion website — update product name to Sevari.

SECTION 08

PostgreSQL Schema Extensions

Serenity's 33-table schema is the foundation. Sevari adds 5 tables. Total: 38 tables.

TABLE	PURPOSE	KEY FIELDS
`sevari_tasks`	All submitted tasks with routing and status	task_id, business_id, target_agent, task_type, subject, body, status, priority, submitted_at, completed_at, source_product ('sevari')
`sevari_agent_runs`	Execution log for every agent invocation	run_id, task_id, agent_name, model_used, tier_used, prompt_tokens, output_tokens, duration_ms, status, output_preview, gate_level, approval_status
`sevari_briefs`	Morning/evening brief history	brief_id, business_id, brief_type ('morning'|'evening'), content, tasks_summarised, approvals_pending, sent_at

`sevari_approvals`	Approval queue with full context	approval_id, task_id, run_id, agent_name, action_type, preview, gate_level, status ('pending'|'approved'|'rejected'), decided_by, decided_at, hive_learned
`sevari_agent_config`	Per-business agent configuration	business_id, agent_name, enabled, gate_level_override, model_override, custom_prompt, schedule_cron, active

All 6 immutable audit tables from Serenity carry forward unchanged. The `source_product` field on `sevari_tasks` and `audit_trail` enables clean separation when Serenity launches customer-facing.

SECTION 09

Pricing — Serenity Tiers

Same four tiers as Serenity. Same prices. Zero revenue share.

TIER	MONTHLY	TASKS/MONTH	AGENTS	FEATURES
Starter	£79	500	BASTION + SENTINEL	Guardian, failsafe chain, basic dashboard, brand voice, PII encryption, immutable audit
Growth	£149	2,000	+ HERALD + VANGUARD	+ HIVE Intelligence, pattern detection, KB gaps, morning briefs, content calendar
Scale	£249	10,000	+ COMPASS + LEDGER	+ Market intelligence, financial tracking, queue monitoring, health dashboard, advanced HIVE
Enterprise	£399	Unlimited	All 7 (+	+ Code generation,

			FORGE)	GitHub integration, full visibility, custom rules, SLA, white-label, 24/7

> **COMPETITIVE PRICING ADVANTAGE**
>
> Competitors charge up to 20% of your revenue OR $500+/month for comparable features. Sevari: fixed price, zero revenue share, all AI inference included (runs locally on the VPS, not metered cloud API calls). The more your business grows, the cheaper Sevari becomes relative to percentage-based competitors.

SECTION 10

n8n Single Canvas Architecture

One workflow. One canvas. Serenity v6.0 as the base (95 nodes) plus Sevari extensions (~35 new nodes). Total: ~130 nodes.

Canvas Zones

ZONE	X POSITION	STATUS	NODES
1: Multi-Input Intake	-300 → 600	INHERITED	IN:API, IN:Form, IN:Webhook, IN:Email, IN:Legacy, IN:Merge, IN:Normalise, IN:Queue + **IN:Task** (new)
2: Guardian Threat Protection	600 → 1800	INHERITED	GRD:Scan, GRD:AiPrep, GRD:PhiScan, GRD:AiEval, GRD:Gate, GRD:Block
3: PII + Credentials	1800 → 3900	INHERITED	SEC:EncryptPII, CRED:*, CORE:Validate
4: Agent	4400 →	NEW	SVR:Detect, SVR:Route

Router	4800		
5: Agent Execution	5000 → 7200	NEW	SVR:FORGE, SVR:HERALD, SVR:VANGUARD, SVR:BASTION, SVR:LEDGER, SVR:COMPASS, SVR:SENTINEL + failsafe merge
6: Core + RAG (Serenity path)	4200 → 7800	INHERITED	CORE:Classify, RAG:*, AI tiers T1-T5 (enquiry path)
7: HIVE Mind + Output	7800 → 10000	INHERITED	ECO:Registry, HIVE:Learn, SEC:RespGuard, OUT:Format, PG writes
8: Operations	0 → 1400, Y:2000+	INHERITED + EXTENDED	All OP:* nodes + SVR:TaskAPI, SVR:AgentStatus, SVR:BriefHistory, SVR:ApprovalQueue
9: Schedulers	0 → 400, Y:4500+	INHERITED + EXTENDED	SCH:* + SVR:SCH:HERALD, SVR:SCH:COMPASS, SVR:SCH:LEDGER, SVR:SCH:SENTINEL, SVR:SCH:Brief
10: Self-Marketing	1600 → 3200, Y:4000+	INHERITED	MKT:* (absorbed into HERALD schedule)

Key Routing Logic

After CORE:Validate passes, a new IF node `SVR:Detect` checks if the incoming data has a `target_agent` field or arrived via the `/sevari/task` webhook. If yes → Sevari agent pipeline. If no → Serenity enquiry pipeline (existing CORE:Classify → RAG → AI tiers path). Both paths converge at HIVE:Learn for shared intelligence.

SECTION 11

Deployment Sequence

PHASE	WHAT	MILESTONE
1	Import Sevari n8n JSON to VPS n8n instance. Create PostgreSQL credentials. Assign to all PG: nodes. Run 38-table SQL schema.	Sevari running on VPS
2	Pull all 6 Ollama models. Test each model responds via `/api/chat` . Run selftest endpoint.	All AI models operational
3	Seed demo business (Hadfield Plumbing BIZ-DEMO-001). Submit test tasks via each agent path. Verify Guardian blocks threats, agents route correctly, HIVE Mind learns.	Live demo operational
4	Configure production domain. Set up reverse proxy (Caddy/nginx) with SSL. Replace ngrok with permanent URL. Update webhook URLs.	Production-accessible
5	Update Operion website: add Sevari product page. Connect onboarding webhook to Sevari. Enable Stripe integration. Deploy legal docs (updated for Sevari).	Global launch
6	HIVE Mind accumulates real-world data. When ready, launch Serenity customer-facing with pre-trained HIVE Mind.	Serenity launch

SECTION 12

VPS Setup Guide

Ollama Models to Pull

```
ollama pull mistral:7b
ollama pull llama3.2:3b
ollama pull gemma2:2b
ollama pull phi3:mini
ollama pull qwen2.5-coder:7b
ollama pull nomic-embed-text
```

Verify Models

```
curl http://localhost:11434/api/chat -d '{
  "model": "mistral:7b",
  "messages": [{"role": "user", "content": "Hello"}],
  "stream": false
}'
```

PostgreSQL Schema

```
psql -U serenity -d serenity_db -f serenity_schema_v6_sevari.sql
```

n8n Import

Import the JSON via n8n UI → Workflows → Import from file. Create PostgreSQL credentials in n8n Settings → Credentials → PostgreSQL. Assign to all `PG:` nodes. Activate the workflow.

> **POST-IMPORT CHECKLIST**
>
> 1. PostgreSQL credentials created and assigned to ALL PG: nodes. 2. All 6 Ollama models pulled and responding. 3. Selftest endpoint returns ALL_PASSED. 4. Demo task submitted to each agent — all 7 produce output. 5. Guardian blocks a test threat payload. 6. Morning brief schedule fires at 07:00. 7. HIVE Mind records patterns from test runs.

END OF DOCUMENT

OPERION SEVARI v1.0 — Operion Automation Systems Ltd — April 2026

operionautomation@gmail.com